

# **Auditable Surge Planning: Network-Aware Causal Inference Meets Prescriptive Optimization**

Sanjay Kumar Patnala & Vedika Sanjeev Lakhanpal

**CML KDD 2025**

# Motivation - why tactical surges matter

- Consumer-goods plants increasingly run short, 30-40 % “surge” runs to rescue service levels when demand spikes or inventory dips.
- Today these surges are triggered by **heuristics** – spreadsheet rules like “boost the SKU with the biggest forecast gap.”
- Result: ad-hoc firefighting → stockouts, premium freight, lost sales.
- We ask: **Can we learn the true impact of a surge and prescribe the optimal set of SKUs to surge next, under real plant constraints?**

# Key Technical Challenges

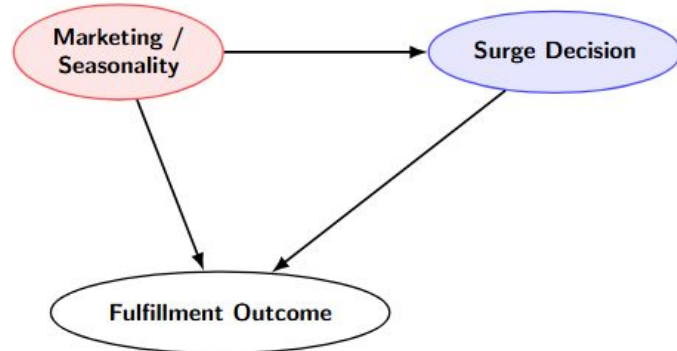
## Resource coupling (network spillover)

- A filler line shared by Sports Drink (A) and Iced Tea (B) means surging A can starve B.

## Confounding (causal attribution)

- Surges coincide with promotions, seasonality, unlogged fixes.
- Naïve before/after comparison confuses surge effect with these shocks.

Need a graph-aware causal framework that separates true surge uplift from confounders and quantifies spillovers on neighbours.



# Related Work

Theme	Key Insights
Production interventions as exogenous shocks	Early works like <i>Gupta et al. (2019)</i> used DiD to estimate surge effects. However, DiD assumes surge timing is unrelated to demand shifts a weak assumption in real plants.
Improved causal methods at coarse granularity	<i>Ahmed et al. (2024)</i> used doubly robust models but at the plant-month level, masking SKU-level variation. <i>Lin et al. (2023)</i> studied overtime spillovers but did not model SKU heterogeneity.
Network interference underexplored in industry	<i>Kim &amp; Hollingsworth (2022)</i> and <i>Vasiliev &amp; Weng (2023)</i> analyzed COVID-19 and retail shocks in networks but lacked fine-grained SKU-level or surge-relevant modeling.
Prescriptive optimization over causal effects	<i>Bertsimas et al. (2020)</i> and <i>Poupart et al. (2022)</i> optimize over treatment effects but ignore estimation uncertainty. Our work incorporates bootstrap quantiles in a chance-constrained knapsack.

# Detecting Surge Episodes

## Rule/Heuristic :

- Compute 3-day moving average output:  $\mu_i(3)(t)$
- Compute 14-day baseline:  $\mu_i(14)(t)$
- Surge starts at day **t** if:  $\mu_i(3)(t) \geq 1.3 \times \mu_i(14)(t)$  and condition holds for  $t, t+1, t+2$ .
- Enforce 10-day cool-down  $\Rightarrow$  17 632 non-overlapping episodes across 2 104 SKUs, spanning Jan–Aug 2023.

Why this matters:

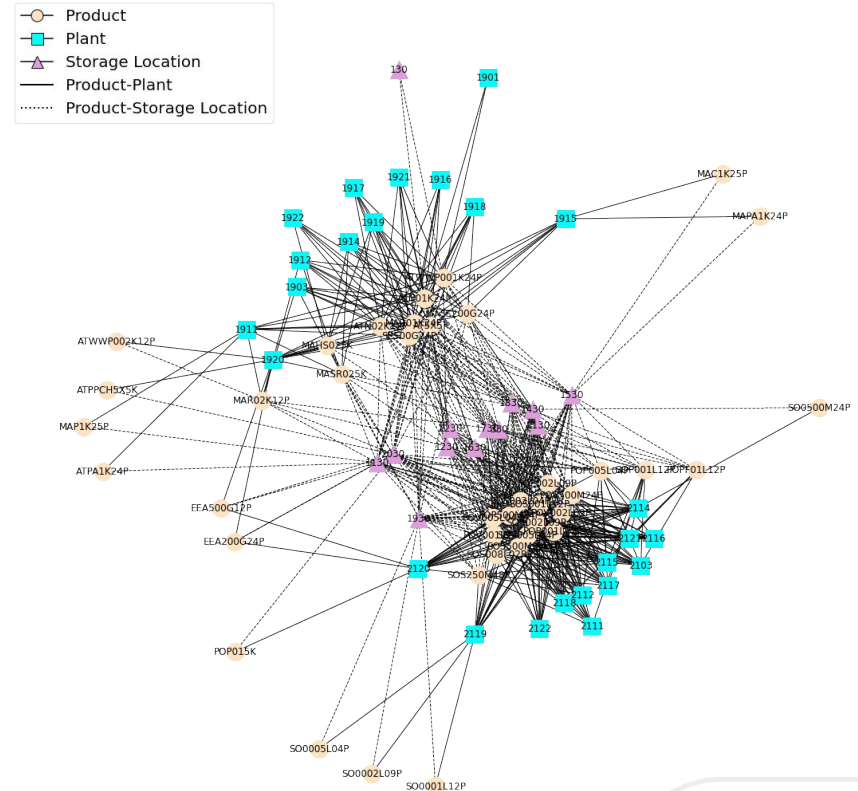
- ✓ Provides an objective treatment label for causal analysis.
- ✓ Matches plant KPI for “tactical surge,” boosting practitioner trust.

# Key contributions

Contribution	Description
Graph-fused propensity model	Smooths treatment probabilities over the SKU resource graph, boosting overlap and keeping coefficients interpretable.
Dynamic marginal structural models	Recovers day-by-day direct surge effects and 7-day cumulative uplift.
Spillover & heterogeneous uplift estimation	Honest uplift forests quantify how much each neighbour benefits or suffers, with per-SKU confidence intervals.
Chance-constrained knapsack optimizer	Selects SKU surge set that maximizes 5th-percentile fulfilment gain under labour / line-hour budgets; solves in under 90 ms.

# Dataset

- Data – SupplyGraph: Open dataset of 2,104 SKUs over 243 days, yielding 17,632 detected surge episodes.
- Each episode has a start day and spans one week of outcome observation. Resource-sharing network: SKUs are nodes connected if they share a production line, primary storage site, or product subgroup.
- The resulting undirected graph has 115,372 edges (average degree ~55), reflecting how capacity and materials link the SKUs. This network encodes potential interference between the SKUs' production.



# Causal Assumptions behind the Pipeline

- **Sequential ignorability**

After conditioning on covariates  $X_{\{it\}}$ , surge assignment  $A_{\{it\}}$  and neighbour exposure  $E_{\{it\}}$  are independent of future outcomes.

- **Partial interference**

Treatment effects propagate **only** along edges in  $G$ ; no long-range spillover beyond the immediate resource-sharing graph.

- **Positivity (overlap)**

Every SKU appears in both surged and non-surged states across 243 days.

- **Control checks**

✓ Post-weight absolute SMD < 0.05 for every feature group.

✓ Only 0.8 % rows have weight > 10 → no extreme extrapolation.



# Graph Based Propensity Estimation

## Model and Graph penalty

$$\Pr(A_{it} = 1 \mid X_{it}) = \sigma(\alpha + \gamma_i + X_{it}^\top \beta)$$

$$\min_{\alpha, \beta, \gamma} - \sum_{i,t} \left[ A_{it} \log p_{it} + (1 - A_{it}) \log(1 - p_{it}) \right] + \lambda \sum_{(i,j) \in E} W_{ij} (\gamma_i - \gamma_j)^2.$$

## Why fuse?

- Nearby SKUs share crews & maintenance schedules → similar propensities.
- Fusion shrinks noisy intercepts, enlarges overlap.

Metric	Plain Logit	+Graph Fusion
Avg  SMD	0.06	<b>0.04</b>
ESS (Effective Sample Size)	910	<b>1,470 (+62%)</b>

# Dynamic Marginal Structural Model (MSM)

Weighted regression ( $h = 0 \dots 7$ )

$$Y_{i,t+h} = \alpha_h + \psi_h A_{it} + \eta_h E_{it} + X_{it}^\top \beta + \varepsilon_{i,h}, \quad h = 0, \dots, 7.$$

## Key findings

- Day 0:  $\psi_0 = -2.6$  pp (congestion)
- Day 3:  $\psi_3 = +1.9$  pp
- Day 7:  $\psi_7 = +5.8$  pp (95 % CI +4.6,+6.9)

Interpretation  $\rightarrow$  backlog clears within 5 days; net uplift after 1 week.

# Heterogeneous Uplift with Honest Forests

## Why heterogeneous?

- 30 % of SKUs show strong overlap → tailor surges where payoff is largest.

## Two learners evaluated

Model	RMSE ↓	R <sup>2</sup> ↑	PICP
DR-Learner (GBM + RF)	1.31	0.28	92%
Honest Uplift Forest	1.18	0.34	94%

Forest captures non-linear interactions → better targeting.

# Quantifying Network Spillovers

## Neighbour model

$$Y_{j,t_k+7} = \alpha + \beta A_{i,t_k} + \eta E_{i,t_k} + X_{j,t_k}^\top \theta + \varepsilon_{jk}$$

## Result

- Average spillover  $\beta = +0.89$  pp (95 % CI +0.48,+1.31)  
→ Coordinated surges can lift fulfilment across shared lines.

## Operational insight

- ✓ Capacity synergy outweighs cannibalisation in this plant network.

# Chance-Constrained Knapsack Optimizer

## Objective

$$\max_{x \in \{0,1\}^{|P|}} \sum_{i \in P} \tau_i x_i \quad \text{s.t.} \quad \sum_{i \in P} c_i x_i \leq K_p$$

- $\tau_i$  = 5th-percentile bootstrap CATE     $c_i \approx 0.84$  h / SKU
- Why 5th-percentile? → Guarantees uplift in worst 5 % scenarios.

## Runtime & optimality

- Greedy = optimal for  $K \leq 5$ ; MILP (CPLEX) < 90 ms for 372 SKUs.
- Daily end-to-end pipeline = 7.4 s on commodity CPU.

Metric	Heuristic	Ours
Mean uplift	+0.8 pp	<b>+5.4 pp</b>
5th-pct uplift	+0.2 pp	<b>+4.7 pp</b>
Worst-case risk	—	<b>−88%</b>

# Prescriptive Results vs. Incumbent Heuristic

Offline replay

Policy	Mean ↑	5th-pct ↑	Std-dev ↓
Volume-rank heuristic	+0.8 pp	+0.2 pp	0.6 pp
Point-estimate knapsack	+6.1 pp	+3.3 pp	1.9 pp
Chance-constrained (ours)	+5.4 pp	+4.7 pp	1.1 pp

## Key takeaways

- **4.9×** higher worst-case uplift than heuristic.
- Trades 0.7 pp mean to slash downside risk by **88 %**.
- Prevents  $\approx 1.9$  M late cases over 8 months.

# Auditability, Robustness, Deployment

## Transparency

- ✓ Propensity weights, MSM coefficients, CATE intervals, solver logs all serialised for domain review.

## Robustness checks

- Placebo surge dates  $\rightarrow$  ATT collapses to 0 (no spurious effect).
- Weight-clip 1  $\leftrightarrow$  99 pct  $\rightarrow$   $< 0.4$  pp drift in ATT.
- Rosenbaum bound: uplift  $> 0$  for  $\Gamma \leq 1.35$  (unmeasured bias).

## Deployment footprint

- Python + scikit-learn + CPLEX; one cron job per plant.
- CPU-only: 7.4 s total per day, fits inside MES scheduler SLA.

# Conclusion

**Graph-aware causal ML** turns noisy surge logs into reliable, SKU-level uplift and spillover estimates (ESS +62 %, SMD < 0.05).

**Risk-aware optimisation** lifts forward-week fulfilment by +5.4 pp on average and protects  $\geq 4.7$  pp in the worst 5 %.  
→ 5× improvement over today's spreadsheet heuristic.

**Glass-box & fast:** full audit trail, sub-second solve times, deployable with existing MES.



**THANK YOU!**